

photo: Książ Castle, Poland

Deep Learning in Python

Introduction to Data Mining

Paweł Kasprowski, PhD, DSc.

Problems to be solved

- Classification
 - assign samples to predefined classes
 - [photo of an animal -> species]
 - [text -> author]
 - [scan-path -> novice/expert]
- Regression
 - calculate a value for each sample
 - [images of the house -> price of the house]
 - [eye image -> gaze coordinates]
- Conversion
 - convert object into another object
 - [image -> text description of the image]
 - [voice recording -> text of the speech]

Task and method

- Task: find a function $Y = f(X)$ where
 - X – **sample** (input)
 - Y – result – class, value (output) – **label**
- Method (learning by example):
 - take some number of samples X with known output Y (**examples**, labeled samples)
 - build the function based on these examples (learning process)
 - use the function to predict the label for unknown samples

Problem of "generality"

- It is relatively easy to create a function that correctly classifies all examples
- But is this function "general" enough?
 - is it able to correctly classify unknown samples?
- The answer is not trivial and that is why we have a lot of different classification methods
- "No free lunch" theorem

Over-fitting

- If the function (model) is optimized for the given data (given examples) it may have poor generality
- This problem is called over-fitting
- Solution is to simplify the model, e.g. stop learning even when it is not perfect for examples

The simplest algorithms

- K-Nearest Neighbors (kNN)
 - a new sample is classified as majority of its K nearest neighboring examples
- Linear Discriminant Analysis (LDA)
 - finds a linear equation separating positive and negative samples
- Naive Bayes
 - the class of a new sample is calculated based on examples using Bayes equation
- Decision Tree
 - the tree is built using examples by searching for the most discriminative features
 - new samples are going through the tree reaching a leaf with the predicted class

More sophisticated algorithms

- Random Forest
 - a set of randomly generated decision trees voting for the final results
- Support Vector Machines (SVM)
 - recalculates examples into another space where they are easier to separate lineary (kernel trick)
- Neural Networks
 - based on neurons organized into layers

Hyperparameters


- Classification algorithms have parameters
 - (so called hyperparameters)
- kNN – number of neighbors
- Random Forest – number of trees
- SVM – kernel function, parameters of the function, learning rate etc.
- Good choice of hyper parameters influences the quality of the model!
- Bad news: deep learning methods have A LOT hyper parameters

Scikit-learn

- Package implementing:
 - many popular classification algorithms
 - many classic methods for data handling
- General use:
 - load the dataset consisting of samples and their labels
 - create a model using one of the possible algorithms (classes)
 - train the model: `model.fit(samples, labels)`
 - use the model to predict classes:
 - `predicted = model.predict(sample)`

Universal example

python classifier



Creating a model

```
def classification(model_name,samples,labels,rangex,rangey):  
models = {  
    "KNN": KNeighborsClassifier(),  
    "LDA": LinearDiscriminantAnalysis(),  
    "NB": GaussianNB(),  
    "TREE":DecisionTreeClassifier(),  
    "RF":RandomForestClassifier(n_estimators=20),  
    "SVM":SVC(gamma='scale'),  
    "PERC":Perceptron(max_iter=2000),  
    "GB":GradientBoostingClassifier()  
}  
model = models.get(model_name)  
...
```

Using the model

```
...
samples = np.array(samples)
labels = np.array(labels)
model.fit(samples, labels)
# build the matrix of results using the model
result = np.zeros([rangex,rangey])
for x in range(rangex):
    for y in range(rangey):
        sample = np.array([x,y])
        result[x][y] = model.predict(sample.reshape(1, -1))
return result
```


Classification of a single sample

- `result = model.predict(samples)`
 - samples must have the dimension:
 - [number of samples, number of attributes]
 - For one sample: [1,2]
- Code version 1:
 - `sample = np.array([x,y])` # dimension [2,]
 - `result[x][y] = model.predict(sample.reshape(1, -1))`
- Code version 2:
 - `sample = np.array([[x,y]])` # double brackets - dimension [1,2]
 - `result[x][y] = model.predict(sample)`

Classification of all samples

- It is possible to classify all samples at once
- Create a list of points to classify:
 `samples = []`
 for x in range(rangex):
 for y in range(rangey):
 `samples.append((x,y))`
- A more pythonic way:
 `samples = [(x,y) for x in range(rangex) for y in range(rangey)]`
- Convert to numpy
 `samples = np.array(samples)`

Classification of all samples at once

- Predictions for all points
`r = model.predict(samples)`
- Create the array of results (shape: [rangex,rangey])
`result = np.zeros([rangex,rangey])`
for i in range(len(samples)):
 `result[samples[i,0],samples[i,1]]=r[i]`
- Is it possible to do it without the "for" loop?

Example for the iris dataset

iris.ipynb

```
import pandas as pd
from sklearn.neighbors import KNeighborsClassifier
data = pd.read_csv('iris.data')
```

```
model = KNeighborsClassifier()
samples = data.values[:,0:4] # first four columns
labels = data.values[:,4] # the fifth column
model.fit(samples, labels)
predicted = model.predict([(5.9, 3.0, 5.1, 1.8)])
print("predicted",predicted)
> predicted ['Iris-virginica']
```

	sl	sw	pl	pw	iris
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa

Evaluating results

- Hyperparameters may be tuned
 - it is important to be able to check if the model is general
- First rule: never check the model using the same data that you used for training (examples)
 - over-fitting!
- The dataset should always be divided into:
 - training set (used to build the model)
 - test set (to check how it works for unknown samples)

Evaluation

- Dividing into training set and test set may be not enough!
 - we optimize the model (by tuning hyper parameters) for the given test set
- The safest way:
 - training set (for learning)
 - validation set (for hyper parameters tuning)
 - test set (for final evaluation)
 - the "holdout" dataset

Sklearn train_test_split

- (trainSamples, testSamples, trainLabels, testLabels) = sklearn.model_selection.train_test_split(samples, labels)
- Two pairs:
 - trainSamples, trainLabels
 - testSamples, testLabels
- Default: 75% train, 25% test
- Parameters:
 - train_size – percent or number of samples
 - test_size – percent or number of samples

iris.ipynb / Train Test Split

Stratification

- Problem: proportion of samples for classes may differ in training and test sets
- For instance: 50-50-50 in the original set
 - training: 40-30-25
 - test: 10-20-25
- Effect: class0 will be treated as more probable!
- Solution: stratified division (maintain distribution)
- Parameter: stratify=labels

Check test samples

iris.ipynb / Check an error

- Checking the test set

```
correct = 0;
```

```
predictedLabels = model.predict(testSamples)
```

```
for i in range(len(testSamples)):
```

```
    print(testLabels[i],"->",predictedLabels[i],end=' ')
```

```
    if(testLabels[i]==predictedLabels[i]):
```

```
        correct = correct + 1; print('OK')
```

```
    else:
```

```
        print('error!!!')
```

```
print("Correct: {} of {} accuracy = {:.2f}"
```

```
    .format(correct,len(testSamples),correct/len(testSamples)))
```

Cross-validation

- Dynamic division into training and test
 - the same examples are sometimes training and sometimes test samples
- 10-fold cross validation:
 - divide the whole dataset into 10 subsets
 - for each subset
 - train the model using the remaining 9 subsets
 - test the results using the chosen subset
 - average the results
- Folds are randomized but may be stratified
 - the same distribution of classes in each fold

Sklearn cross_validate

- ***iris.ipynb / Cross validation***
 - `sklearn.model_selection.cross_validate(model, samples, labels, cv=<number of folds>)`
- Returns for each fold:
 - `fit_time`
 - `score_time`
 - `test_score`
- Many parameters to check other metrics!

Measures

- Accuracy – the number of correctly classified samples to all samples
 - a problem with unbalanced sets
 - if 90% of test samples is positive the blind classifier achieves 90% accuracy
- Precision – the cost of false positives
 - not belonging to class predicted as belonging
- Recall – the cost of false negatives
 - belonging to class predicted as not belonging
- F1-Score – combination of precision and recall
- Cohen's kappa – compares prediction with random prediction

Measures interpretation

- Example: We have built the model that diagnoses COVID-19 based on a genetic test
- For every eye movement sample the model returns:
 - ***H*** - if a person is healthy
 - ***S*** - when a person is sick
- We test the model using
 - 18 samples from sick people (***S***)
 - 82 samples from healthy people (***H***)
- We achieve 82% accuracy
 - is it bad or good?
 - it depends: we must examine a confusion matrix

Confusion matrix

Predicted as >> Real class	H	S
H	81	1
S	17	1

P	N
TP	FN
FP	TN

Accuracy 0.82

H:
precision: 0.83
recall: 0.99
F1-score: 0.90

S:
precision: 0.50
recall: 0.06
F1-score: 0.1

Cohen's kappa: 0.07

Measures

Predicted as >> Real class	H	S
H	80	2
S	15	3

P	N
TP	FN
FP	TN

Accuracy 0.83

H:
precision: 0.84
recall: 0.98
F1-score: 0.90

S:
precision: 0.60
recall: 0.17
F1-score: 0.26

Cohen's kappa: 0.20

Measures

Predicted as >> Real class	H	S
H	66	16
S	2	16

P	N
TP	FN
FP	TN

Accuracy 0.82

H:
precision: 0.97
recall: 0.80
F1-score: 0.88

S:
precision: 0.50
recall: 0.89
F1-score: 0.64

Cohen's kappa: 0.53

Measures

Predicted as >> Real class	H	S
H	64	18
S	0	18

P	N
TP	FN
FP	TN

Accuracy 0.82

H:
precision: 1.00
recall: 0.78
F1-score: 0.88

S:
precision: 0.50
recall: 1.00
F1-score: 0.67

Cohen's kappa: 0.56

Example for three classes

Confusion matrix

```
[[24 3 9]  
 [ 3 24 24]  
 [ 2 15 49]]
```

class	precision	recall	f1-score	support
0	0.83	0.67	0.74	36
1	0.57	0.47	0.52	51
2	0.60	0.74	0.66	66
micro avg	0.63	0.63	0.63	153
macro avg	0.67	0.63	0.64	153
weighted avg	0.64	0.63	0.63	153

Accuracy: 0.63

Scikit learn measures

- Useful functions from *sklearn.metrics*:
 - `confusion_matrix(labels, results)`
 - `classification_report(labels, results)`
 - `accuracy_score(labels, results)`
 - `precision_score(labels, results)`
 - `recall_score(labels, results)`
 - `cohen_kappa_score(labels, results)`
 - ...

Measures calculation

iris.ipynb / Measures calculation

```
from sklearn.metrics import classification_report, confusion_matrix,  
    accuracy_score,cohen_kappa_score
```

```
predictedLabels = model.predict(testSamples)  
print(confusion_matrix(testLabels, predictedLabels))  
print(classification_report(testLabels, predictedLabels))  
accuracy = accuracy_score(testLabels, predictedLabels)  
print("Accuracy: {:.2f}".format(accuracy))  
c_kappa = cohen_kappa_score(testLabels, predictedLabels)  
print("Cohen's Kappa: {:.2f}".format(c_kappa))
```

Adding own scorers to cross validation

iris.ipynb / Adding own scorers to cross validation

```
from sklearn.metrics import make_scorer, precision_score, recall_score
```

```
p_scorer = make_scorer(precision_score, average='micro')
```

```
r_scorer = make_scorer(recall_score, average='micro')
```

```
my_scorer = {'precision': p_scorer, 'recall': r_scorer}
```

```
sklearn.model_selection.cross_validate(model, samples, labels, cv=5,  
    scoring=my_scorer)
```

Summary

- Accuracy is often not enough
 - especially for unbalanced datasets
- Cohen's kappa is the reliable way to evaluate results
 - but has some drawbacks as well!
- The evaluation depends on the target we want to achieve
 - False Acceptance Rate ($FAR=0$)
 - no incorrectly classified healthy people
 - False Rejection Rate ($FRR=0$)
 - no incorrectly classified sick people

Regression

- The model does not search for one of N classes but for a value
 - E.g. price of the house
- Examples:
 - Linear Regression
 - Support Vector Regression
 - Decision Tree Regression
 - Neural Network Regression

Evaluation of regression

- Mean Absolute Error
- Mean Squared Error
 - penalizes predictions differing greatly (outliers)
- Root Mean Squared Error
 - comparable to real values
- Coefficient of Determination (R^2)
 - value 0-1
 - negative value if worse than random

Example: wine quality dataset

- Every wine has 11 numerical features
 - 1 - fixed acidity
 - 2 - volatile acidity
 - 3 - citric acid
 - 4 - residual sugar
 - 5 - chlorides
 - 6 - free sulfur dioxide
 - 7 - total sulfur dioxide
 - 8 - density
 - 9 - pH
 - 10 - sulphates
 - 11 - alcohol
- The result: wine quality in 0-10 scale
- Dataset: 4898 wines



Wine dataset analysis

wine.ipynb

- Simple classification
- Train-test split
- Measures calculation – confusion matrix
- Problem:
 - quality is an ordinal value not just a class
 - results should be calculated as a distance from the correct one
 - it is not bad when wine with quality 3 is classified as 4
 - it is a big problem when it is classified as 9

Regression measures

- Errors:

```
from sklearn.metrics import mean_squared_error, r2_score  
print('MSE=', mean_squared_error(testLabels, predictedLabels))  
print('R2=', r2_score(testLabels, predictedLabels))
```

- Results:

- MSE = 0.99
- R2 = -0.27

Linear Regression

- Using Linear Regression:

```
from sklearn import linear_model  
model = linear_model.LinearRegression()  
model.fit(trainSamples, trainLabels)  
predictedLabels = model.predict(testSamples)
```
- The result for each sample is a float!
- Errors:
 - $MSE = 0.53$
 - $R^2 = 0.31$

Errors per class estimation

- Mean errors per class


```
errors = np.abs(modelResults-testLabels)
```

```
for i in range(3,10):
```

```
    print('class',i,' avg error=',errors[np.where(testLabels==i)].mean())
```

- Results

- class 3 avg error= 2.33
- class 4 avg error= 1.31
- class 5 avg error= 0.61
- class 6 avg error= 0.31
- class 7 avg error= 0.78
- class 8 avg error= 1.71
- class 9 avg error= 2.42



Problem: imbalanced classes

- Some qualities are rare, some are common

- Solution: Calculate weight for each class:

```
from sklearn.utils import class_weight
class_weights = class_weight.compute_class_weight('balanced',
                                                  classes=np.unique(trainLabels),y=trainLabels)
weights = np.ones([len(trainLabels)]) # initialize all weights to 1
for i, label in enumerate(trainLabels):
    weights[i] *= class_weights[int(label-3)] # the first label is 3!
```

- Using weights:

```
model.fit(trainSamples, trainLabels, sample_weight=weights)
```

Regression with weights

- Mean errors per class


```
errors = np.abs(modelResults-testLabels)
```

```
for i in range(3,10):
```

```
    print('class',i,' avg error=',errors[np.where(testLabels==i)].mean())
```

- Results

- class 3 avg error= 2.06
- class 4 avg error= 1.02
- class 5 avg error= 0.66
- class 6 avg error= 0.70
- class 7 avg error= 0.72
- class 8 avg error= 1.39
- class 9 avg error= 1.48



Results

- Distribution is more equal for the weighted model
- But the error is higher (both MSE and R2)
- Question:
 - what do we really need?
 - what is our objective?
- Possibilities:
 - minimize absolute error
 - minimize MSE/RMSE
 - minimize error for high quality wines
 - ...

Feature selection

- Not every feature is valuable
 - adds some knowlegde about a class
 - e.g. bottle's color is *probably* not significant for wine quality
- Not important features may spoil classification!
- There are plethora of algoritms that aim at selecting only the relevant features from the dataset
 - topic for a new lecture (or even a new subject)
- Sklearn has a very convenient class to solve the problem

SelectKBest

- Code:

```
from sklearn.feature_selection.univariate_selection import  
    SelectKBest  
  
newSamples = SelectKBest(<algorithm>, k=<number>)  
    .fit_transform(samples, labels)
```

- Available algorithms:

(see: https://scikit-learn.org/stable/modules/generated/sklearn.feature_selection.SelectKBest.html)

f_classif, mutual_info_classif, chi2, f_regression,
mutual_info_regression, SelectPercentile, SelectFpr,
SelectFdr, SelectFwe, GenericUnivariateSelect


Example

- Select 5 best features using f_regression algorithm:

```
from sklearn.feature_selection.univariate_selection import SelectKBest  
print("Samples before", samples.shape)  
newSamples = SelectKBest(sklearn.feature_selection.f_regression, k=5)  
    .fit_transform(samples, labels)  
print("Samples after", newSamples.shape)  
samples = newSamples
```
- Output:
 - Samples before (4898, 10)
 - Samples after (4898, 5)
- Task for you: check if it works better!

Simple linear regression

- Find the best linear approximation of points



Example: Linear Regression

- Given a training set of pairs (X,Y) find coefficients of a linear function

$$y = a * x + b$$

- Error (loss) function Mean Squared Error

$$E = \frac{1}{n} \sum_i^n (y_i - \hat{y}_i)^2 = \frac{1}{n} \sum_i^n (y_i - (ax_i + b))^2$$

- Goal: minimize this function

Linear Regression using Gradient Descent

- Gradient Descent algorithm:
 - initialize randomly a and b
 - repeat:
 - calculate error (loss)
 - change a and b in such way that error is smaller
 - until error is small enough or it is not decreasing
- Question: how to find the direction of change for a and b ?

Gradient Descent

- Gradient dE/da – a function describing how E changes when a changes

$$\frac{\delta E}{\delta a} = \frac{-2}{n} \sum_i^n (\hat{y}_i - y_i) x_i$$

- Gradient dE/db - function describing how E changes when b changes

$$\frac{\delta E}{\delta b} = \frac{-2}{n} \sum_i^n (\hat{y}_i - y_i)$$

Gradient Descent algorithm

- for each step (iteration)
 - calculate error E
 - calculate gradients dE/da and dE/db
 - $a = a - \text{learning rate} * dE/da$
 - $b = b - \text{learning rate} * dE/db$
- Every iteration gives better values of a and b
 - error (**loss**) is smaller
- Learning rate controls how fast the parameters are updated
 - too fast – we may pass the minimum
 - too slow – we may wait forever

Analytical solution

- For linear regression the best parameters may be calculated using equations:

$$a = \frac{\sum_i^N (x_i - \bar{x})(y_i - \bar{y})}{\sum_i^N (x_i - \bar{x})^2}$$


$$b = \bar{y} - a\bar{x}$$

- but the presented example shows how Gradient Descent algorithm works (and we will use this algorithm later!)

Example

linear_descent.ipynb

- Simple code calculating linear regression using gradient descent for a set of points




Observations

- If we change initial values of a and b we have different results
 - b is reluctant to change
 - kind of "vanishing gradient" problem
- (1) Change the learning rate for b:
 - $nb = b - L * 100 * D_b$ # Update b
- (2) Change the number of iterations
 - epochs = 100 (1000?)
- The results are better!

Learning rate

- Change to:
 - $L = 0.0004$ # The learning rate
- Results jump over the solution
 - too big step!



Summary

- There are many different classification algorithms
- There are many ways to tune these algorithms
- There are many measures to estimate the quality of classification/regression
- So what is so special about the deep learning methods?
 - Wait for the following sections!

Deep Learning with Python

Next lecture:
Neural Networks in Keras/Tensorflow

Introduction to Data Mining

Paweł Kasprowski, PhD, DSc.

